

LEZ. 13/11/2024 part 2

Elective exercise using Go and RPC

Corso di Sistemi Distribuiti e Cloud Computing
A.A. 2024/25

Valeria Cardellini

Laurea Magistrale in Ingegneria Informatica

Elective exercise using Go and RPC

- Realize a distributed application that solves the sorting problem using **MapReduce**

The beauty of MapReduce is that any programmer can understand it, and its power comes from being able to harness thousands of computers behind that simple interface.

David Patterson

- Requirements:
 - Use either **Go and RPC** or **Go and gRPC**
 - Organize properly your code into separate files
 - 1 student per team (2 students only if also the optional part is implemented)

esempio

dataset di
ingresso

15
6
12
2
6
125
3
1
8

MASTER

dividere
dataset
in chunk

MAP₁

MAP₂

MAP₃

+
svolge il
ruolo di
client

faccio ordinamento
in modo crescente

6
12
15

2
6
125

1
3
8

output
intermedio

odi di
reduce

RED₁

RED₂

↓
svolge il
ruolo di
server

valori che
comprende

merge 1-12

13-125

merge 15
125

file 1
di
output

F₁

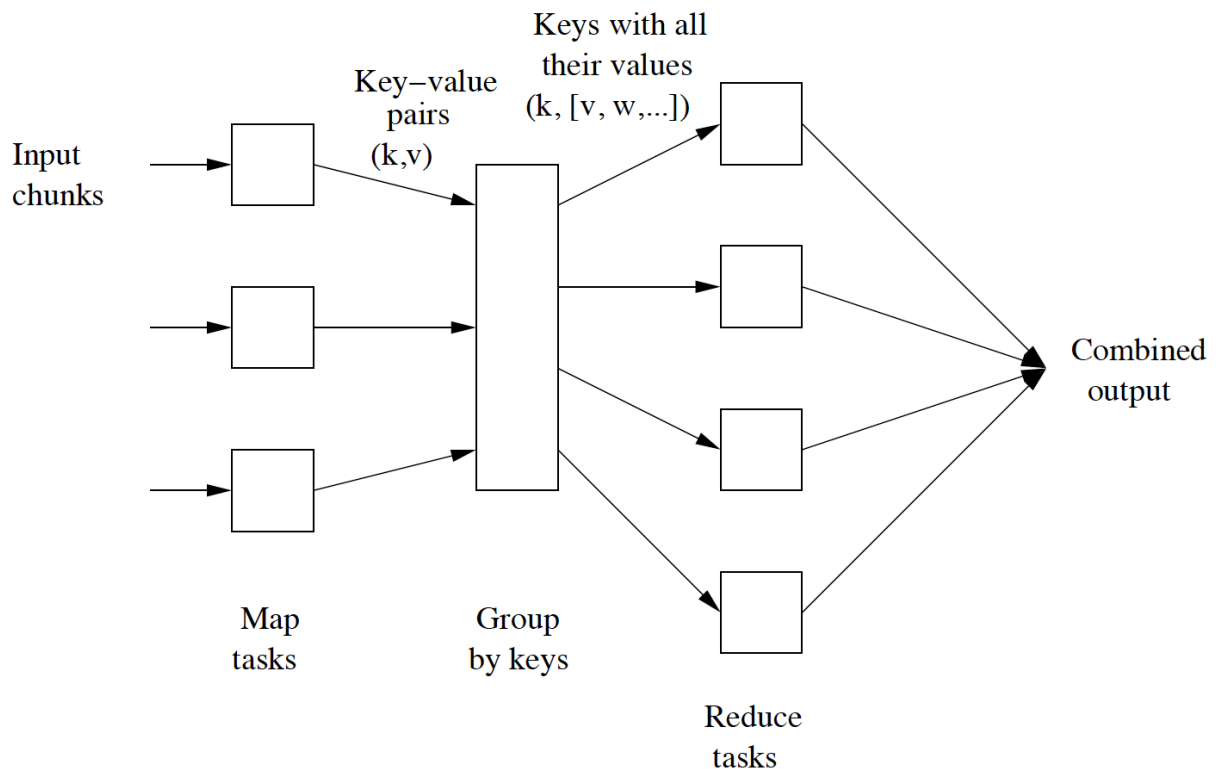
F₂

file 2
di
output

reduce → chiamato Terasort

salva il
risultato
finale in
un file

MapReduce paradigm




Valeria Cardellini - SDCC 2024/25

2

Overview on architecture

- Exploit master-worker architecture
 - Distribute work among workers using RPC or gRPC
- Master assigns **map and reduce tasks** to workers
- Workers execute map and reduce tasks
- Do not consider failures of master and workers
 - Set of workers is known and does not change during computation; no worker fails
 - Master does not fail



Valeria Cardellini - SDCC 2024/25

3

Overview on architecture

- 3 phases
 - **Map**: sort the chunk received from master
 - **Shuffle**: map tasks send their intermediate results to reduce by partitioning result
 - **Reduce**: merge intermediate results received from map
- Distribute work among parallel workers
- Need a **synchronization** point (i.e., barrier) between map and reduce tasks
 - No reduce task can start until all the map tasks have finished their processing
- Each reduce task writes its results on a file

Some simplifying assumptions

- Master and workers not fail during computation → *nessun fallimento durante computaz.*
- The set of workers is known to the ster when the system starts and is defined into a configuration file
- The ports are defined into a configuration file → *indirizzo e porta*

*progetto : almeno 2 nodi di reduce
(ma dovrebbe funzionare anche
con un singolo nodo)*

Optional

→ per chi fa progetto
in due

- Sample input data to optimize partitioning among reduce task
- You can containerized your distributed application
 - To build the image, see Go official image hub.docker.com/_/golang
 - A Docker container per application component and then use Docker Compose to orchestrate the multiple containers on your computer

Delivery

- When
 - By **December 13, 2024**
- What
 - Your code, including a README with instructions to run it
 - Optional: very short report describing the architecture of the distributed application and its main features
- How
 - By email
 - Use as mail subject: **[SDCC] consegna esercizio in Go**

istruzioni minimali per
esecuzione dell'app

Titolo email